

# Assignment 3: Group 11

## FlightDelays dataset

Goal: Accurately predict whether or not a new flight will be delayed (outcome 1 = delayed and 0 = on time).

Data: 2001 record with 5 predictors. All flights from the Washington, DC area into the New York City area during January 2004.

### Variables in the dataset:

CRS\_DEP\_TIME: Computer Reservation Systems departure Time

CARRIER: Airlines

DEP\_TIME: Actual Departure time

DEST: Destination Airport

DISTANCE: Distance travelled by flight

FL\_DATE: Flight date

FL\_NUM: Flight Number

ORIGIN: Origin Airport

Weather: Weather conditions

DAY\_WEEK: Weekdays starting from 1 to 7, 1 = Monday, 2 = Tuesday, 3 = Wednesday..., 7 = Sunday

DAY\_OF\_MONTH: Day of a month

TAIL\_NUM: Tail number

Flight Status: Status of the flight that is ontime or delayed.

```
In [50]: #Importing Libraries

import numpy as np
import math
import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import accuracy_score, roc_curve, auc
import matplotlib.pyplot as plt
from dmba import regressionSummary, classificationSummary
from dmba import liftChart, gainsChart
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import export_graphviz
import graphviz
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.linear_model import LogisticRegression
```

```
In [2]: # Reading the dataset
flight_df = pd.read_csv("FlightDelays.csv")
flight_df.sample(10)
```

Out[2]:

	CRS_DEP_TIME	CARRIER	DEP_TIME	DEST	DISTANCE	FL_DATE	FL_NUM	ORIGIN
1985	1455	OH	1455	JFK	184	1/29/2004	5935	BV
1333	630	DH	627	EWR	213	1/19/2004	7305	IA
2005	1230	DL	1230	LGA	214	1/29/2004	1752	DC
1043	1400	MQ	1400	LGA	214	1/15/2004	4966	DC
1040	800	MQ	800	LGA	214	1/15/2004	4954	DC
1049	630	US	627	LGA	214	1/15/2004	1479	DC
19	1830	MQ	1822	JFK	213	01/01/2004	4784	DC
1126	1000	US	958	LGA	214	1/16/2004	2166	DC
1563	800	US	754	LGA	214	1/22/2004	2162	DC
1566	1100	US	1055	LGA	214	1/22/2004	2168	DC

```
In [3]: #Displaying all columns
flight_df.columns
```

Out[3]: Index(['CRS\_DEP\_TIME', 'CARRIER', 'DEP\_TIME', 'DEST', 'DISTANCE', 'FL\_DATE',  
'FL\_NUM', 'ORIGIN', 'Weather', 'DAY\_WEEK', 'DAY\_OF\_MONTH', 'TAIL\_NUM',  
'Flight Status'],  
dtype='object')

```
In [4]: flight_df.dtypes
```

Out[4]: CRS\_DEP\_TIME int64  
CARRIER object  
DEP\_TIME int64  
DEST object  
DISTANCE int64  
FL\_DATE object  
FL\_NUM int64  
ORIGIN object  
Weather int64  
DAY\_WEEK int64  
DAY\_OF\_MONTH int64  
TAIL\_NUM object  
Flight Status object  
dtype: object

## Target Variable - Flight Status: ontime or delayed

```
In [5]: flight_df['Flight Status'].unique()
```

```
Out[5]: array(['ontime', 'delayed'], dtype=object)
```

### A. Which variables are quantitative/numerical? Which are ordinal? Which are nominal?

From the variables listed above:

#### Quantitative/Numerical Variables:

CRS\_DEP\_TIME (Integer)  
DEP\_TIME (Integer)  
DISTANCE (Integer)  
FL\_NUM (Integer)  
Weather (Integer)  
DAY\_WEEK (Integer)  
DAY\_OF\_MONTH (Integer)

#### Categorical Variables:

Nominal Variables:

CARRIER (String)  
DEST (String)  
FL\_DATE (String)  
ORIGIN (String)  
TAIL\_NUM (String)  
Flight Status (Boolean)

In [6]: `flight_df.describe()`

Out[6]:

	CRS_DEP_TIME	DEP_TIME	DISTANCE	FL_NUM	Weather	DAY_WEEK	DAY_OF_MONTH
<b>count</b>	2201.000000	2201.000000	2201.000000	2201.000000	2201.000000	2201.000000	2201.000000
<b>mean</b>	1371.938664	1369.298955	211.871422	3815.086324	0.014539	3.905498	15.000000
<b>std</b>	432.697149	442.462754	13.316815	2409.750224	0.119725	1.903149	1.903149
<b>min</b>	600.000000	10.000000	169.000000	746.000000	0.000000	1.000000	1.000000
<b>25%</b>	1000.000000	1004.000000	213.000000	2156.000000	0.000000	2.000000	2.000000
<b>50%</b>	1455.000000	1450.000000	214.000000	2385.000000	0.000000	4.000000	4.000000
<b>75%</b>	1710.000000	1709.000000	214.000000	6155.000000	0.000000	5.000000	5.000000
<b>max</b>	2130.000000	2330.000000	229.000000	7924.000000	1.000000	7.000000	7.000000

In [7]: `#Creating a subset dataframe with only numerical variables for statistical  
subset_flight_df = flight_df[["CRS_DEP_TIME", "DEP_TIME", "DISTANCE", "FL_NUM", "Weather", "DAY_WEEK", "DAY_OF_MONTH"]]  
subset_flight_df.head()`

Out[7]:

	CRS_DEP_TIME	DEP_TIME	DISTANCE	FL_NUM	Weather	DAY_WEEK	DAY_OF_MONTH
<b>0</b>	1455	1455	184	5935	0	4	1
<b>1</b>	1640	1640	213	6155	0	4	1
<b>2</b>	1245	1245	229	7208	0	4	1
<b>3</b>	1715	1709	229	7215	0	4	1
<b>4</b>	1039	1035	229	7792	0	4	1

## A. Data Preprocessing

### Do some statistics on the variables

In [8]:  *#Compute mean, standard deviation, min, max, median, length, and missing v*

```
pd.DataFrame({'mean': subset_flight_df.mean(),
              'Sd': subset_flight_df.std(),
              'min': subset_flight_df.min(),
              'max': subset_flight_df.max(),
              'median': subset_flight_df.median(),
              'length': len(subset_flight_df),
              'miss.val': subset_flight_df.isnull().sum(),
              })
```

Out[8]:

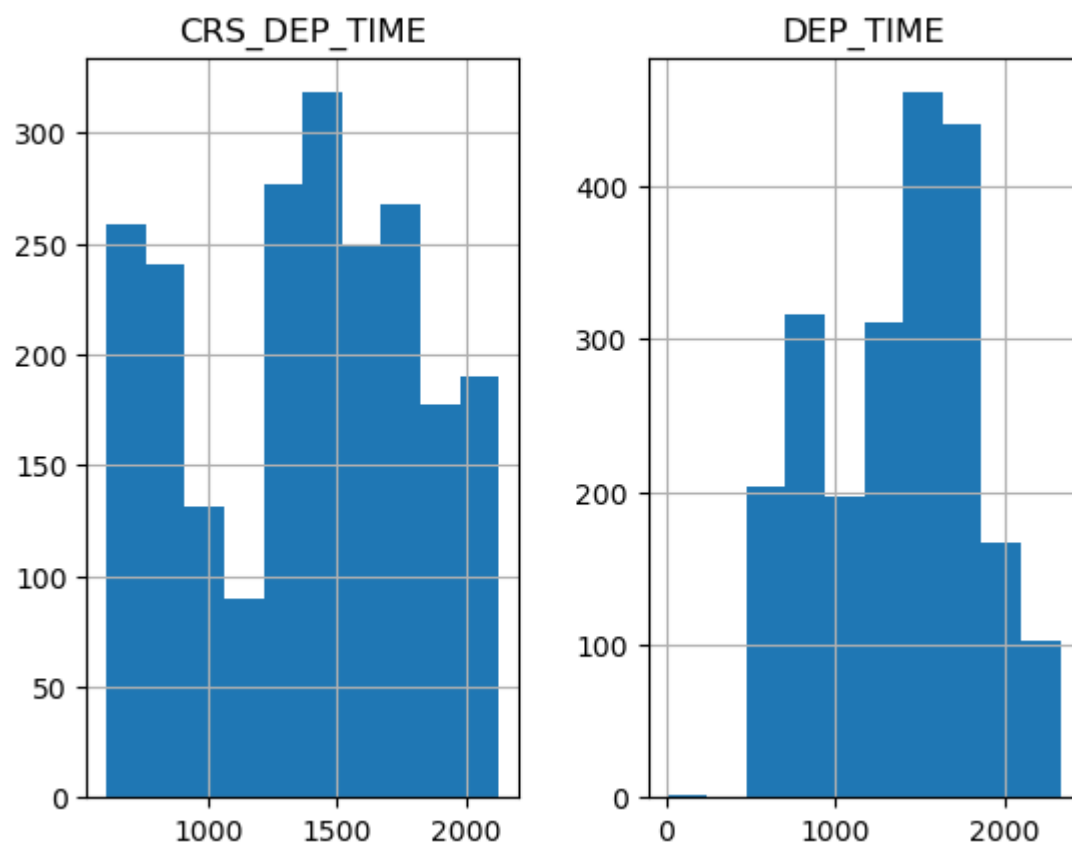
	mean	Sd	min	max	median	length	miss.val
<b>CRS_DEP_TIME</b>	1371.938664	432.697149	600	2130	1455.0	2201	0
<b>DEP_TIME</b>	1369.298955	442.462754	10	2330	1450.0	2201	0
<b>DISTANCE</b>	211.871422	13.316815	169	229	214.0	2201	0
<b>FL_NUM</b>	3815.086324	2409.750224	746	7924	2385.0	2201	0
<b>Weather</b>	0.014539	0.119725	0	1	0.0	2201	0
<b>DAY_WEEK</b>	3.905498	1.903149	1	7	4.0	2201	0
<b>DAY_OF_MONTH</b>	16.024989	8.677390	1	31	16.0	2201	0

## C. Using histogram and summary statistics

```
In [9]: fig, axes = plt.subplots(1, 2)

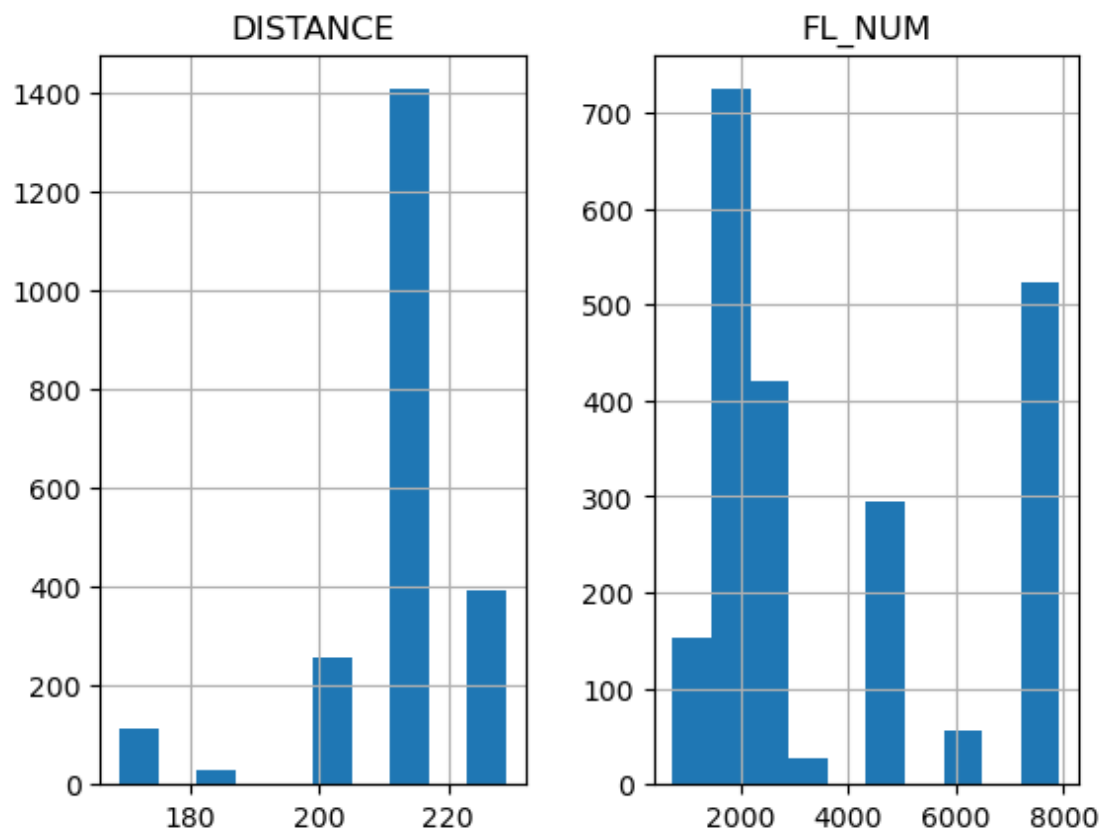
subset_flight_df.hist('CRS_DEP_TIME', ax=axes[0])
subset_flight_df.hist('DEP_TIME', ax=axes[1])

plt.show()
```



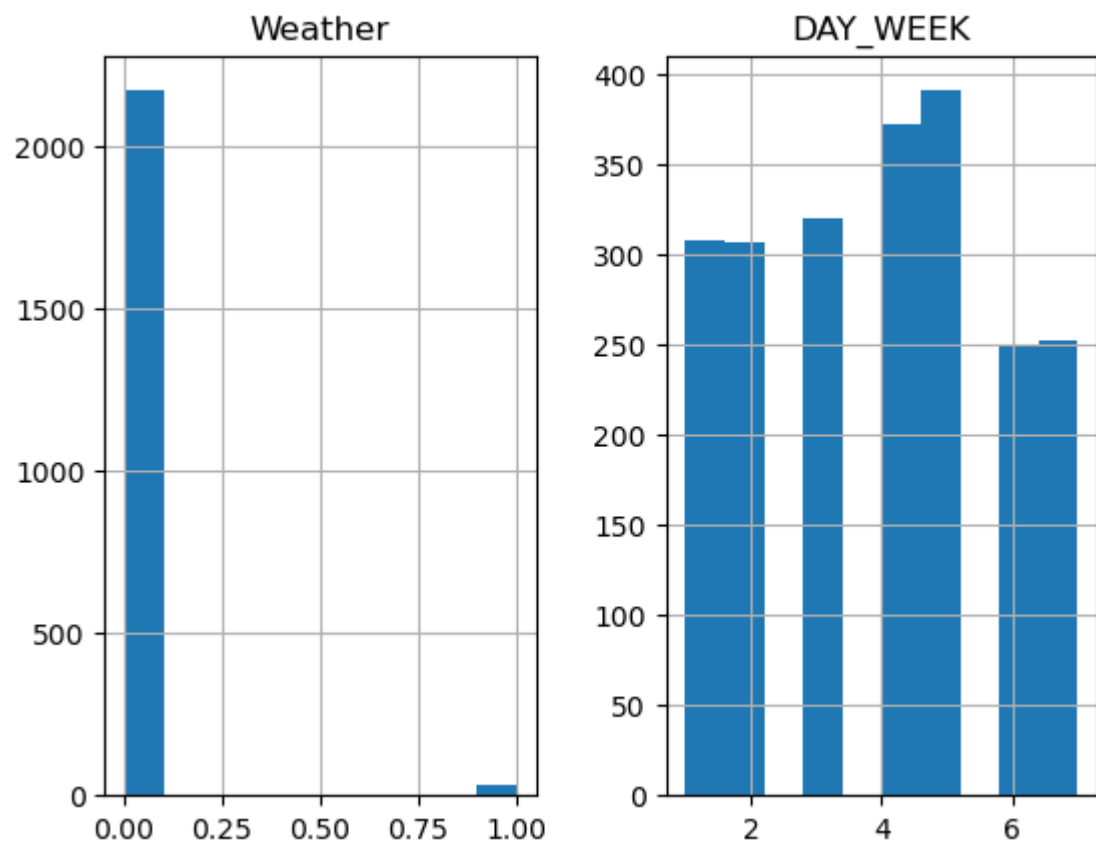
```
In [10]: fig, axes = plt.subplots(1, 2)

subset_flight_df.hist('DISTANCE', ax=axes[0])
subset_flight_df.hist('FL_NUM', ax=axes[1])
plt.show()
```



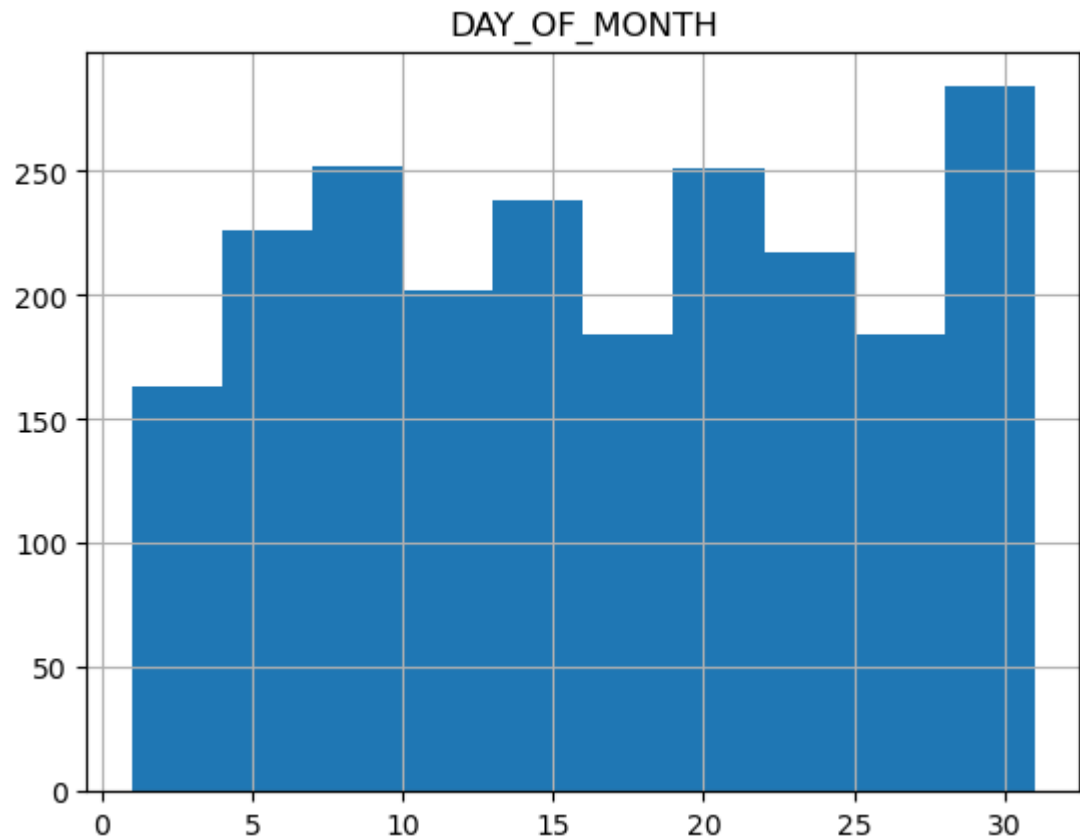
```
In [11]: fig, axes = plt.subplots(1, 2)

subset_flight_df.hist('Weather', ax=axes[0])
subset_flight_df.hist('DAY_WEEK', ax=axes[1])
plt.show()
```





```
In [12]: subset_flight_df.hist('DAY_OF_MONTH')  
plt.show()
```

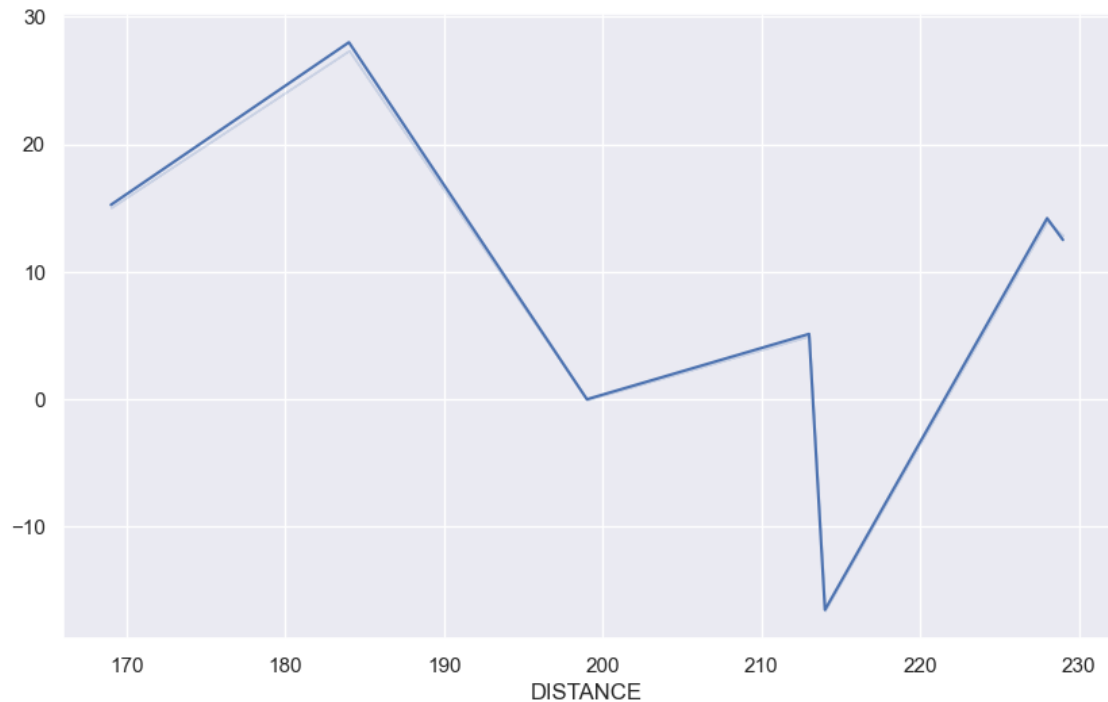


**Using different plots to compare between variables.**

**1. What is the relation of Delay and Distance of a flight of our Dataset?**

```
In [13]: sns.set(rc={"figure.figsize":(10, 6)})  
sns.lineplot(x=subset_flight_df['DISTANCE'], y=(subset_flight_df['DEP_TIME
```

```
Out[13]: <AxesSubplot:xlabel='DISTANCE'>
```



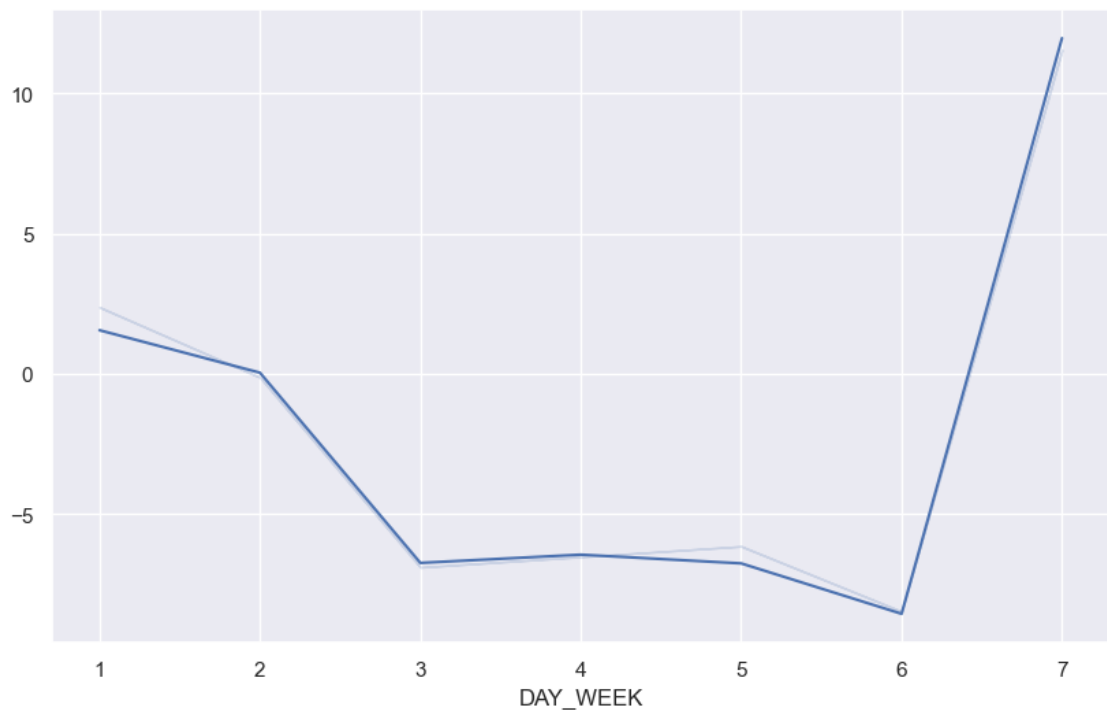
The figure above shows the delay is higher for the flights at smaller distance rather than the flights at longer distance.

This makes sense as generally short distance flights are generally the cheaper flights, which run a jam packed schedule which creates a threat of delay even as there is a very cool down period between each journey.

## 2. Which day of the week witnesses the highest delay.

```
In [14]: sns.lineplot(x=subset_flight_df['DAY_WEEK'],y=(subset_flight_df['DEP_TIME']
```

```
Out[14]: <AxesSubplot:xlabel='DAY_WEEK'>
```



We can see according to our data most of the delay happens on 7th day of the week, i.e. Sunday

## Correlation between Quantitative variables

```
In [15]: # Checking correlation between variables using corr() function
```

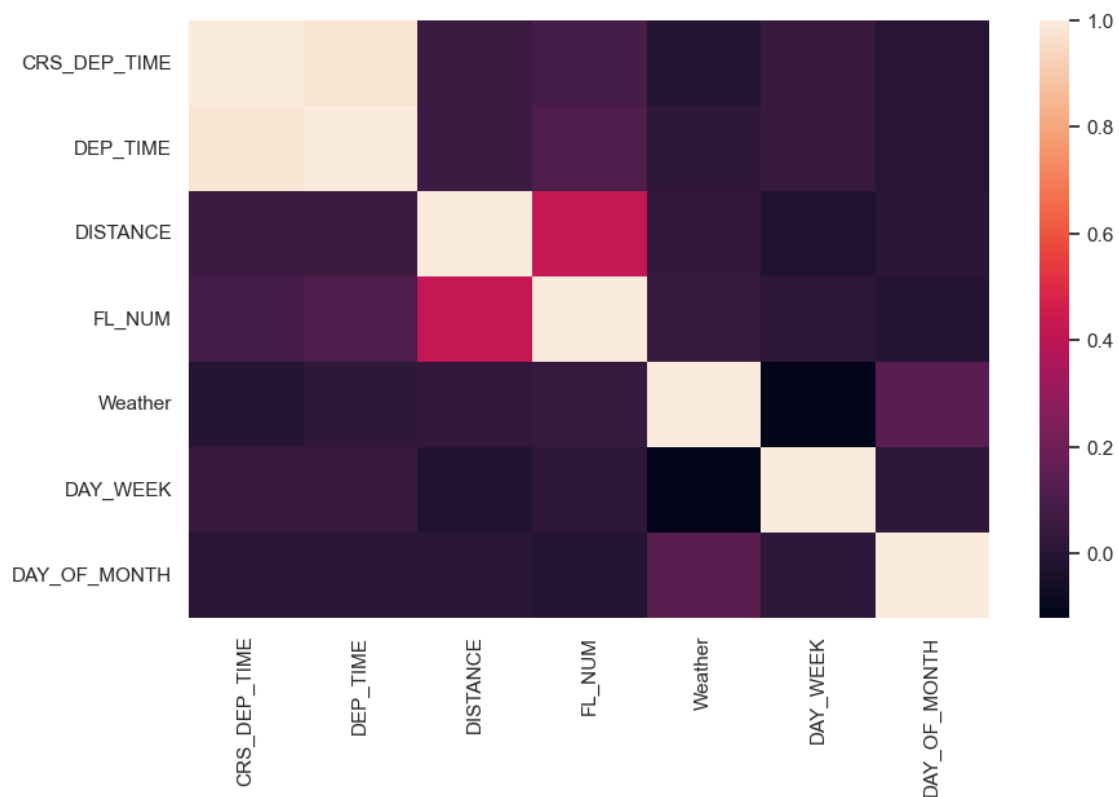
```
subset_flight_df.corr().round(2)
```

```
Out[15]:
```

	CRS_DEP_TIME	DEP_TIME	DISTANCE	FL_NUM	Weather	DAY_WEEK	DAY_OF_MONTH
CRS_DEP_TIME	1.00	0.98	0.06	0.09	-0.01	0.05	0.00
DEP_TIME	0.98	1.00	0.06	0.11	0.02	0.05	0.00
DISTANCE	0.06	0.06	1.00	0.42	0.03	-0.02	0.01
FL_NUM	0.09	0.11	0.42	1.00	0.04	0.02	-0.01
Weather	-0.01	0.02	0.03	0.04	1.00	-0.12	0.14
DAY_WEEK	0.05	0.05	-0.02	0.02	-0.12	1.00	0.02
DAY_OF_MONTH	0.00	0.00	0.01	-0.01	0.14	0.02	1.00

```
In [16]: # Using seaborn  
# Simple heatmap of correlations (without values)  
  
import seaborn as sns  
corr = subset_flight_df.corr().round(2)  
sns.heatmap(corr)
```

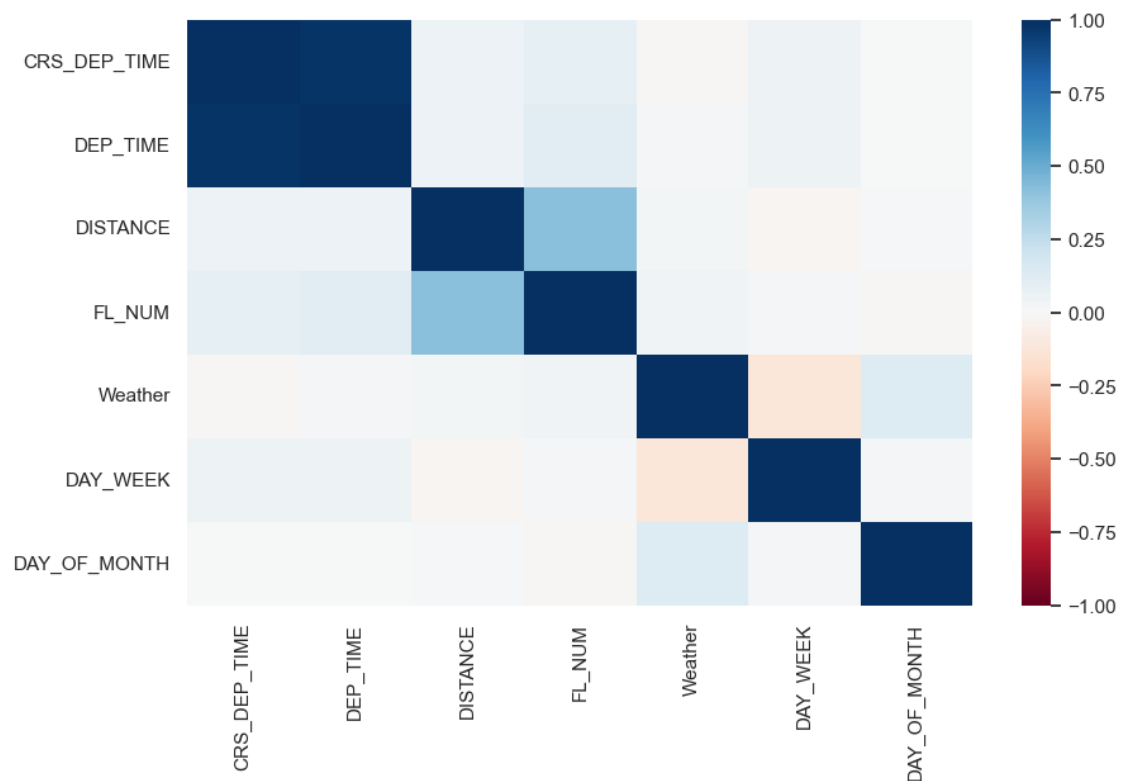
Out[16]: <AxesSubplot:>



In [17]: `# Change to divergent scale and fix the range`

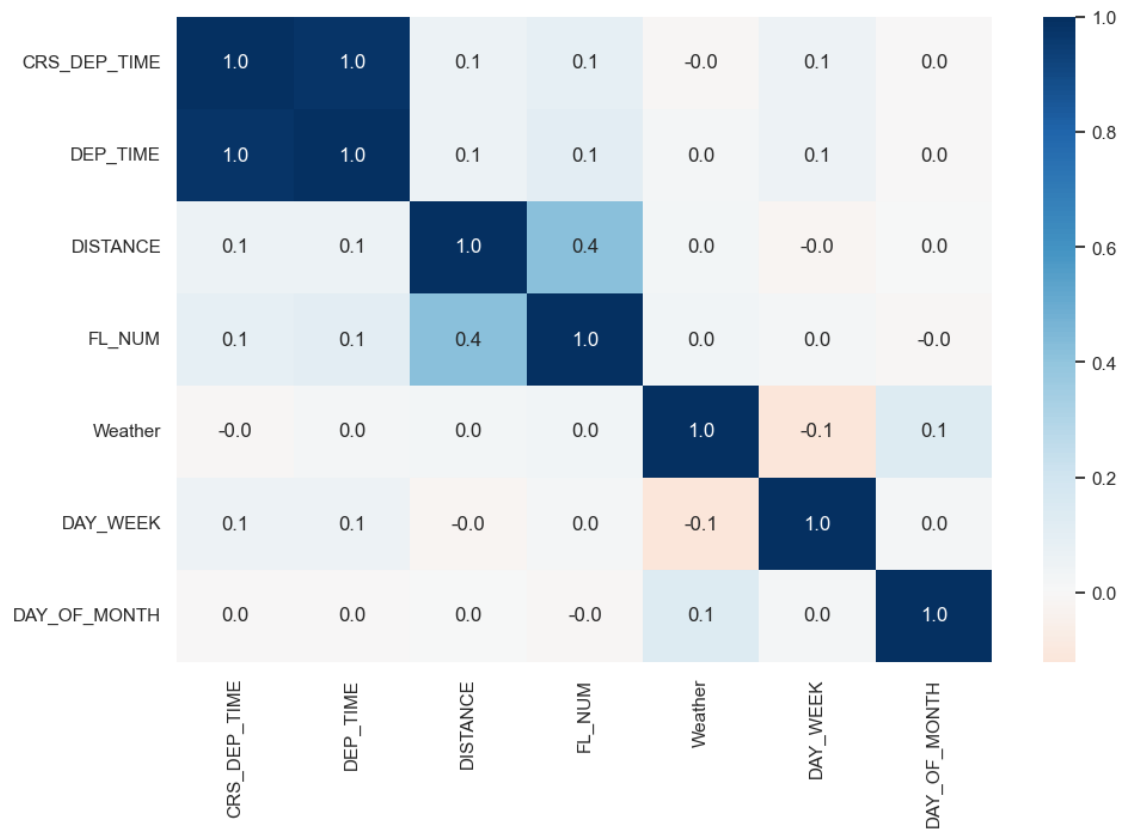
```
sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns, vmin=-1, vmax=1, cmap="RdBu")
```

Out[17]: `<AxesSubplot:>`



```
In [18]: # Include information about values (example demonstrates how to control the
fig, ax = plt.subplots()
fig.set_size_inches(11, 7)
sns.heatmap(corr, annot=True, fmt=".1f", cmap="RdBu", center=0, ax=ax)
```

Out[18]: <AxesSubplot:>



## Correlation between variables:

1. 'CRS\_DEP\_TIME' and 'DEP\_TIME' are strongly correlated with a positive correlation of 1.0 (100% correlation)

## PCA (Principle Component Analysis)

```
In [19]: #Applying PCA

from sklearn.decomposition import PCA
pcs = PCA(n_components=2)
pcs.fit(subset_flight_df)
```

Out[19]: PCA(n\_components=2)

```
In [20]: # To explain the PCA use:
         pcs.explained_variance_ratio_
```

```
Out[20]: array([0.9387127 , 0.06074699])
```

### Removing the Correlated Variables

```
In [21]: flight_df = flight_df.drop(columns=['CRS_DEP_TIME'])
         flight_df.columns
```

```
Out[21]: Index(['CARRIER', 'DEP_TIME', 'DEST', 'DISTANCE', 'FL_DATE', 'FL_NUM',
               'ORIGIN', 'Weather', 'DAY_WEEK', 'DAY_OF_MONTH', 'TAIL_NUM',
               'Flight Status'],
              dtype='object')
```

```
In [22]: flight_df.shape
```

```
Out[22]: (2201, 12)
```

```
In [23]: flight_df.sample(10)
```

```
Out[23]:
```

	CARRIER	DEP_TIME	DEST	DISTANCE	FL_DATE	FL_NUM	ORIGIN	Weather	DAY
1301	MQ	1559	LGA	214	1/19/2004	4970	DCA	0	
1127	US	1059	LGA	214	1/16/2004	2168	DCA	0	
1449	DL	931	LGA	214	1/21/2004	1746	DCA	0	
329	DL	1629	LGA	214	01/06/2004	1760	DCA	0	
2013	DL	2031	LGA	214	1/29/2004	1768	DCA	0	
1659	RU	1032	EWR	169	1/23/2004	2303	BWI	0	
1313	US	1353	LGA	214	1/19/2004	2174	DCA	0	
287	CO	1846	EWR	199	01/05/2004	814	DCA	0	
2065	RU	857	EWR	199	1/29/2004	2582	DCA	0	
426	MQ	1352	LGA	214	01/07/2004	4966	DCA	0	

```
In [24]: #Split the data into training (60%) and testing (40%)

predictors = ['CARRIER', 'DEP_TIME', 'DEST', 'DISTANCE', 'FL_DATE', 'FL_NUM',
              'DAY_OF_MONTH', 'TAIL_NUM']
outcome = 'Flight Status'

X = pd.get_dummies(flight_df[predictors], drop_first=True)
y = flight_df[outcome]
classes = ['ontime', 'delayed']

# split into training and validation
X_train, X_valid, y_train, y_valid = train_test_split(X, y, test_size=0.4,
```

```
In [25]: X_train.head(10)
```

```
Out[25]:
```

	DEP_TIME	DISTANCE	FL_NUM	Weather	DAY_WEEK	DAY_OF_MONTH	CARRIER_DI
1215	2110	229	7684	0	7	18	
1476	659	214	2160	0	3	21	
1897	858	214	2164	0	2	27	
83	1258	214	2172	0	5	2	
1172	1509	213	746	0	6	17	
1720	1252	199	808	0	6	24	
1816	929	214	1746	0	1	26	
1544	1831	214	1764	0	4	22	
2065	857	199	2582	0	4	29	
427	1457	214	4968	0	3	7	

10 rows × 595 columns

```
In [26]: X_train.columns
```

```
Out[26]: Index(['DEP_TIME', 'DISTANCE', 'FL_NUM', 'Weather', 'DAY_WEEK', 'DAY_OF_M
ONTH',
               'CARRIER_DH', 'CARRIER_DL', 'CARRIER_MQ', 'CARRIER_OH',
               ...
               'TAIL_NUM_N970DL', 'TAIL_NUM_N973CA', 'TAIL_NUM_N974DL',
               'TAIL_NUM_N975CA', 'TAIL_NUM_N983CA', 'TAIL_NUM_N986DL',
               'TAIL_NUM_N987DL', 'TAIL_NUM_N994DL', 'TAIL_NUM_N995CA',
               'TAIL_NUM_N997DL'],
              dtype='object', length=595)
```

```
In [27]: X_train.shape
```

```
Out[27]: (1320, 595)
```



## B. K-Nearest Neighbors and Naive Bayes Classifier comparison

### Naive Bayes Algorithm

```
In [28]: ▶ # run naïve Bayes
flight_nb = MultinomialNB(alpha=0.01)
flight_nb.fit(X_train, y_train)

# predict probabilities (Shows the belonging probabilities of each record
predProb_train = flight_nb.predict_proba(X_train)
predProb_valid = flight_nb.predict_proba(X_valid)

# predict class membership (shows the class instead of probability by sele
y_valid_pred = flight_nb.predict(X_valid)
y_train_pred = flight_nb.predict(X_train)
```

```
In [29]: ▶ # Use the model to predict a new data

df = pd.concat([pd.DataFrame({'actual': y_valid, 'predicted': y_valid_pred
df.head(10)
```

Out[29]:

	actual	predicted
1276	ontime	delayed
1446	ontime	ontime
335	ontime	delayed
1458	ontime	ontime
2038	ontime	ontime
1314	ontime	ontime
389	ontime	delayed
1639	ontime	delayed
2004	ontime	ontime
403	ontime	ontime

```
In [30]: ▶ from sklearn.metrics import accuracy_score, precision_score, recall_score,
# Calculate accuracy, precision, recall, and F1-score on the test data
accuracy = accuracy_score(y_valid, y_valid_pred)

precision = precision_score(y_valid, y_valid_pred, average='weighted')
recall = recall_score(y_valid, y_valid_pred, average='weighted')
f1 = f1_score(y_valid, y_valid_pred, average='weighted')
```

```
In [31]: ▶ print("Accuracy:", accuracy)
          print("Precision:", precision)
          print("Recall:", recall)
          print("F1-score:", f1)
```

```
Accuracy: 0.5845629965947786
Precision: 0.7189401801847995
Recall: 0.5845629965947786
F1-score: 0.6282279026455564
```

## K-Nearest Neighbor (KNN ) Algorithm

```
In [32]: ▶ from sklearn.neighbors import KNeighborsClassifier
          from sklearn.metrics import confusion_matrix
```

```
In [33]: ▶ knn = KNeighborsClassifier(n_neighbors=3)
```

```
In [34]: ▶ knn.fit(X_train, y_train)
```

```
Out[34]: KNeighborsClassifier(n_neighbors=3)
```

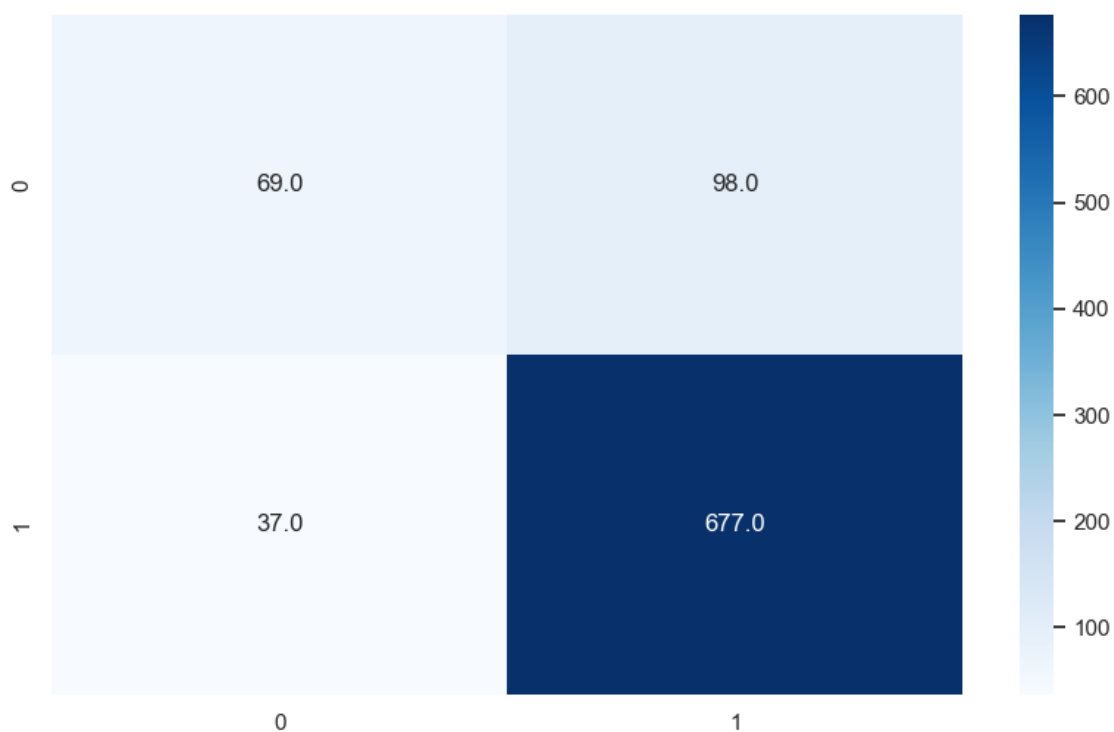
```
In [35]: ▶ y_pred_knn = knn.predict(X_valid)
```

```
C:\Users\sidha\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.
  mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```



```
In [39]: ▶ sns.heatmap(cm_knn,annot=True,cmap = 'Blues', fmt = '0.1f')
```

```
Out[39]: <AxesSubplot:>
```



### Naive Bayes: Confusion Matrix

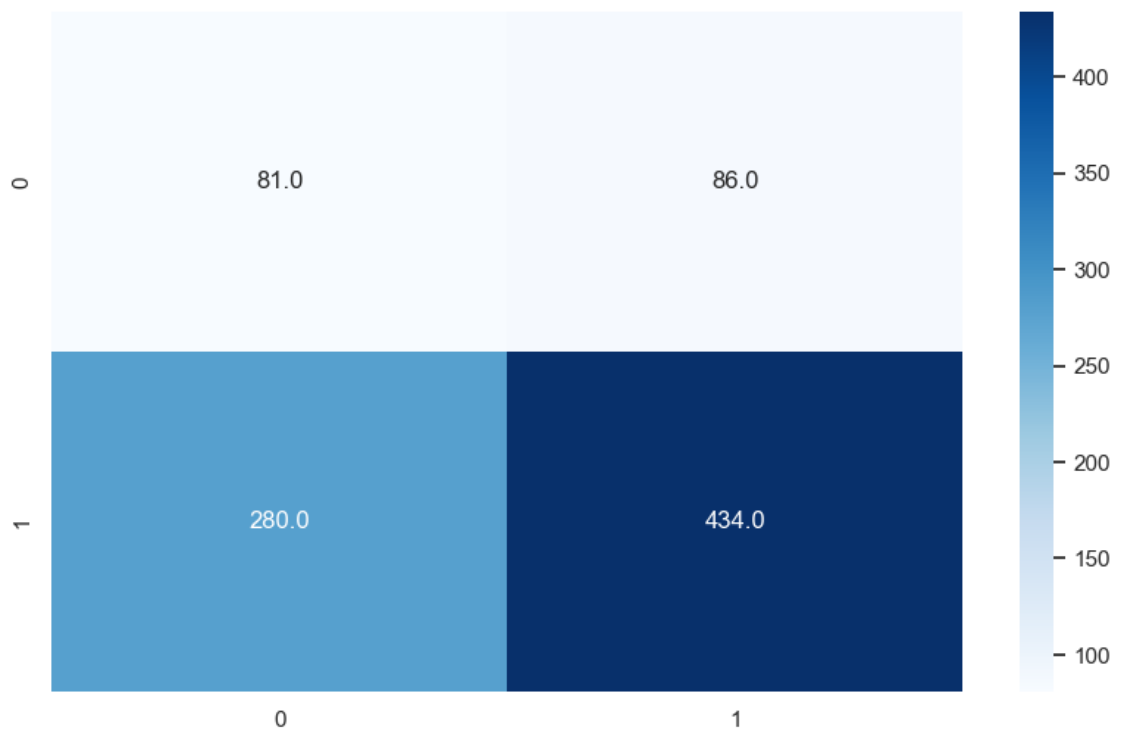
```
In [40]: ▶ cm_nb= confusion_matrix(y_valid,y_valid_pred)
```

```
In [41]: ▶ cm_nb
```

```
Out[41]: array([[ 81,  86],  
               [280, 434]], dtype=int64)
```

```
In [42]: sns.heatmap(cm_nb,annot=True,cmap = 'Blues', fmt = '0.1f')
```

```
Out[42]: <AxesSubplot:>
```



```
In [43]: Accuracy_Knn = (69+677)/(69+98+37+677)
Precision_knn = 69/(69+37)
Recall_knn = 69/(69+98)
F1_Score_knn = 2*Precision_knn*Recall_knn/(Precision_knn+Recall_knn)
```

```
In [44]: Accuracy_nb = (81+434)/(81+86+280+434)
Precision_nb = 81/(81+280)
Recall_nb = 81/(81+86)
F1_Score_nb = 2*Precision_nb*Recall_nb/(Precision_nb+Recall_nb)
```

```
In [45]: acc_score = {"Parameter": ['Accuracy', 'Precision', 'Recall', 'F1-Score'], "Naive Bayes": [Accuracy_nb, Precision_nb, Recall_nb, F1_Score_nb], "KNN": [Accuracy_Knn, Precision_knn, Recall_knn, F1_Score_knn]}
```

```
In [46]: acc_score = pd.DataFrame(acc_score)
```

```
In [47]: acc_score
```

```
Out[47]:
```

	Parameter	Naive Bayes	KNN
0	Accuracy	0.584563	0.846765
1	Precision	0.224377	0.650943
2	Recall	0.485030	0.413174
3	F1-Score	0.306818	0.505495

## Conclusion:

We ran Naive Bayes and K-Nearest Neighbor (KNN) algorithm on Delayed Flights dataset and predicted values for a test dataset for both. After analyzing and comparing the above accuracy measure results, we can say that KNN algorithm gives far better result than Naive Bayes model on our dataset.

1. KNN model gives 84.68 % accurate prediction.
2. Naive Bayes model gives 58.46 % accurate prediction.

## C. Proposing extra algorithm(s)

### 1. CART Algorithm

```
In [55]: ▶ # Create the decision tree classifier and fit the data
flight_dt = DecisionTreeClassifier(max_depth=5)
flight_dt.fit(X_train, y_train)

# Predict the test data using the decision tree and calculate the accuracy
y_pred_dt = dt.predict(X_valid)
accuracy_dt = accuracy_score(y_valid, y_pred_dt)
print("Decision Tree Accuracy: ", accuracy_dt)
```

Decision Tree Accuracy: 0.8229284903518729

```

In [56]: # Fit a decision tree model
clf = DecisionTreeClassifier(max_depth=5)
clf.fit(X_train, y_train)

# Plot the decision tree
plt.figure(figsize=(20,10))
plot_tree(clf, filled=True, rounded=True, feature_names=X_train.columns)
plt.show()

# Generate Graphviz format file
dot_data = export_graphviz(clf, out_file=None,
                           feature_names=X_train.columns,
                           class_names=['ontime', 'delayed'],
                           filled=True, rounded=True,
                           special_characters=True)

# Convert Graphviz format to image format
graph = graphviz.Source(dot_data)
graph.format = 'png'

# Save the image to a file
graph.render('decision_tree')

```



Out[56]: 'decision\_tree.png'

## 2. Logistic Regression

```
In [57]: ▶ # Create the logistic regression classifier and fit the data
flight_lr = LogisticRegression()
flight_lr.fit(X_train, y_train)

# Predict the test data using logistic regression and calculate the accuracy
y_pred_lr = lr.predict(X_valid)
accuracy_lr = accuracy_score(y_valid, y_pred_lr)
print("Logistic Regression Accuracy: ", accuracy_lr)
```

Logistic Regression Accuracy: 0.8195232690124858

C:\Users\sidha\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

## Result:

After comparing the all above algorithms on the FlightDelay dataset, we can conclude that KNN is best fit model with highest accuracy.

1. KNN model gives 84.68 % accurate prediction.
2. Naive Bayes model gives 58.46 % accurate prediction.
3. CART model gives 82.29 % accurate prediction.
4. Logistic Regression model gives 81.95 % accurate prediction.